

PLAIN-ENGLISH EXPLAINER

How My AI & Tech News Platform Works

No jargon — this is written so you can explain it to anyone who asks "so how does your site actually work?"

Prepared: July 20, 2026
Project: AI & Tech News

1. What the Platform Actually Does

Think of it like a smart newspaper front page that never sleeps. It doesn't write news itself — it watches 36 real news websites (TechCrunch, The Verge, BBC, NYT, OpenAI's own blog, and so on), notices when they publish something new, picks out the most important stories, writes a short 2-3 line summary of each one, and shows it on my homepage. Clicking a story sends you to the real article on the real news site — I never copy the full article onto my own site.

The one rule everything follows: never store or show the full article. Only a headline, a short summary, and a link back to the original site. That's what keeps this legal and fair to the publishers.

2. How We Actually Get the News (RSS, explained simply)

Almost every news website secretly publishes a second, machine-readable version of its homepage called an **RSS feed**. RSS stands for "**Really Simple Syndication**" — syndication just means "sharing the same content in many places." You've probably never seen it because it's not meant for humans to browse; it's meant for other computer programs to read automatically.

An RSS feed is basically a simple, constantly-updating list that looks like this for every article: *title, short preview text, link to the full article, and the time it was published*. That's it. No ads, no images baked in, no full article text — just the bare facts of "here's what we just published."

- 1 Every 60 seconds**, my platform quietly visits all 36 of these RSS lists at once (like checking 36 different mailboxes at the same time).
- 2** For each site, it reads whatever new headlines are sitting in that list right now.
- 3** It cleans each headline up — strips away leftover website code, ignores anything older than 7 days, and figures out a topic (like "AI Models" or "Security") and which companies it mentions, just by scanning the words in the title.
- 4** If two sites report basically the same story with a very similar title, only one copy is kept — so the homepage doesn't show the same news three times.

So where does the RSS link for each site actually come from?

This part is **not automatic** — the platform doesn't go hunting the internet for feeds on its own. Every one of the 36 RSS links was found once, by hand, while building the platform, and typed into one settings file. Here's how a person actually finds a site's RSS link:

- A Guess the common address.** Most news sites put their feed at a predictable spot — usually the homepage address plus `/feed`, `/rss`, or `/rss.xml` at the end. For example, TechCrunch's is simply `techcrunch.com/feed/`. This guess works surprisingly often.

- B Look for the little orange icon or an "RSS" link.** Many sites still put a small icon (usually orange, with a dot and two curved lines — the universal RSS symbol) somewhere on their homepage or footer that links straight to the feed.
- C Peek at the page's hidden code.** Every browser can show a webpage's raw underlying code ("view page source"). Sites that offer RSS almost always include one invisible line in that code announcing where their feed lives, even if there's no visible icon or link anywhere on the actual page.
- D Just search for it.** Searching "[site name] RSS feed" almost always turns up the answer immediately, since this is such a common, well-documented thing for websites to have.

Once found, that link is saved permanently in the platform's one settings file (`src/config/sources.js`) — it's checked once by a human, then reused automatically forever after, every 60 seconds, without anyone needing to look it up again.

Does every news site have an RSS feed? (No — and that matters)

No, not all of them. RSS was extremely common in the 2000s-2010s, but it's not a legal requirement or a universal standard every site must offer. Some newer or more locked-down websites have quietly shut their public RSS feeds down over the years — often to push readers toward their own app or newsletter instead. A handful of sites also only offer RSS to paying subscribers, not the public.

This is actually a big part of *why* these specific 36 sources were chosen — every one of them was checked and confirmed to still have a working, public, free RSS feed before being added. Sites that don't offer one simply can't be included this way; the platform intentionally does not try to scrape or copy a site's homepage directly, since that's far more fragile (breaks the moment a site redesigns its page) and much greyer ground legally than reading a feed the site deliberately published for exactly this purpose.

One-line way to explain it: "RSS only works for sites that choose to publish one — so my source list is really 'the 36 major tech and AI outlets that both matter and still offer a public feed,' not literally every tech site in the world."

3. The 36 News Sources — Who They Are & Why They Made the List

These aren't random blogs pulled off the internet. Every one of the 36 is a real, established publisher with actual reporters and editors, or — in a few cases — the company itself posting its own official news. Here's the full list, grouped the same way the code groups them.

Major Tech Press — 18 sources

The big, well-known technology newsrooms most people would already recognize.

Source	Website	What it is
TechCrunch	techcrunch.com	Leading Silicon Valley tech & startup news outlet
The Verge	theverge.com	Major tech, gadgets & culture newsroom (Vox Media)
Ars Technica	arstechnica.com	Long-running, deeply technical tech journalism
Wired	wired.com	Long-established tech & science magazine (Condé Nast)
MIT Technology Review	technologyreview.com	Published by MIT — academic-grade tech reporting
Engadget	engadget.com	Consumer gadgets & tech news, since 2004
ZDNET	zdnet.com	Enterprise & consumer tech news, decades-old outlet
The Register	theregister.com	UK enterprise/IT-focused tech news
CNET	cnet.com	One of the oldest mainstream tech news sites
Gizmodo	gizmodo.com	Gadgets, science & tech culture newsroom
TechRadar	techradar.com	Consumer tech news & reviews (Future plc)
The Next Web	thenextweb.com	European tech & startup news outlet
Digital Trends	digitaltrends.com	Consumer tech news & buying guides
TechSpot	techspot.com	Hardware & computing-focused tech news
Tom's Hardware	tomshardware.com	PC hardware & components news, since 1996
9to5Google	9to5google.com	Dedicated Google/Android news desk
9to5Mac	9to5mac.com	Dedicated Apple news desk
Android Authority	androidauthority.com	Dedicated Android/mobile news desk

Mainstream Press, Technology Desks — 4 sources

Global general-news organizations, but only pulling their dedicated technology section.

Source	Website	What it is
--------	---------	------------

BBC Technology	bbc.com/news/technology	The BBC's technology desk — UK public broadcaster
The Guardian Tech	theguardian.com/technology	The Guardian newspaper's technology section
NYT Technology	nytimes.com/section/technology	The New York Times' technology section
CNBC Technology	cnbc.com/technology	CNBC's technology & business news desk

AI-Focused Publications & Official Company Blogs — 9 sources

Dedicated AI news outlets, plus a few sources that are the company's own official announcements — about as "straight from the horse's mouth" as news gets.

Source	Website	What it is
VentureBeat	venturebeat.com	Tech/business news, AI-category feed specifically
AI News	artificialintelligence-news.com	Dedicated AI industry news outlet
MarkTechPost	marktechpost.com	AI & machine learning research news
The Decoder	the-decoder.com	Dedicated AI news outlet
Google AI Blog	blog.google/technology/ai	Official — Google's own AI announcements
OpenAI	openai.com/news	Official — OpenAI's own announcements
NVIDIA Blog	blogs.nvidia.com	Official — Nvidia's own announcements
AWS ML Blog	aws.amazon.com/blogs/machine-learning	Official — Amazon's AI/ML team announcements
IEEE Spectrum	spectrum.ieee.org	Magazine of IEEE, the world's largest engineering body

Security News — 3 sources

Source	Website	What it is
Bleeping Computer	bleepingcomputer.com	Cybersecurity & breach news, widely cited by researchers
The Hacker News	thehackernews.com	One of the most-read cybersecurity news sites
Krebs on Security	krebsonsecurity.com	Run by Brian Krebs, a well-known investigative security journalist

India Tech & Startups — 2 sources

Source	Website	What it is
Economic Times Tech	economictimes.indiatimes.com/tech	Technology desk of India's largest business newspaper

Inc42	inc42.com	Dedicated India startup & tech ecosystem news outlet
-------	-----------	--

Why these are genuine, and not random blogs

- **All 36 are outlets that already exist and report independently of my platform.** I'm not writing anything myself — I only forward what these newsrooms already published under their own name and reputation.
- **Every card links straight back to them.** If a story ever looked wrong or made-up, it's immediately obvious which real publisher it came from, because you land on their site, not mine.
- **Several are the company's own mouth, not a third party.** OpenAI, NVIDIA, Google AI, and AWS ML are official blogs — when they post, it's a primary source by definition, not someone reporting on someone else.
- **Most have existed for 10-25+ years** (Wired, CNET, TechCrunch, The Register, IEEE, Krebs on Security) with real editorial standards and reputations they have to protect — they're not anonymous or disposable websites.

Why these specific 36 were chosen

The goal was **broad, trustworthy coverage of the whole tech & AI world** without leaning on any single outlet's bias — a mix of: big general tech newsrooms, established mainstream newspapers' tech desks, AI-specialist outlets, official company announcements, dedicated security researchers, and an India-specific angle. Each source also gets an internal "authority" score from 5 to 9 (bigger, more established names like BBC, NYT, TechCrunch, The Verge, and MIT Technology Review score 8-9) which quietly gives their stories a bit more weight when the ranking system (see Section 4) is deciding what counts as the top story — without ever excluding the smaller ones.

Can we add more sources later?

Yes, easily. All 36 sources live in one simple list in the code (`src/config/sources.js`). Adding a new one is just adding four things: the outlet's name, its RSS feed link, its homepage link, and an authority score from roughly 5 to 9. Nothing else in the platform needs to change — the fetching, ranking, topic-tagging, and summarizing all automatically apply to any new source the moment it's added. The only real step is checking beforehand that the new site's RSS feed actually loads properly, so a broken link doesn't sit there silently failing.

4. How the "Top Story" (Big Headline) Gets Picked

The big story at the very top of the homepage isn't chosen by a person, and it isn't just "whatever is newest." Every single story gets a **score** — a number that represents how important it looks — and whichever recent story has the highest score wins the top spot. Think of it like a point system:

What the story is about	Points added
Contains the word "breaking"	+30
A launch, release, or big announcement	+18
A company buying/merging with another	+16
Funding, an IPO, or a big investment	+14
Mentions a major AI model (GPT, Claude, Gemini...)	+14
A dollar amount in the billions	+12
Words like "breakthrough," "record," "first-ever"	+12
Government/regulation/legal news	+12
A hack, data breach, or security scare	+12
Mentions a big-name company (Google, Apple, Nvidia...)	+8
Has a photo attached	+6

A story can rack up several of these at once — e.g. "OpenAI raises \$10 billion in new funding" would score points for the funding keyword, the dollar-amount, and the OpenAI mention, all stacked together. On top of that, each source's own authority score (5-9, from Section 3) is added in too, so a story from a bigger, more established outlet gets a small natural edge.

Two more rules that shape the score

- ↓ **Old news fades out on its own.** A story's score automatically shrinks the older it gets — it loses half its "freshness value" roughly every 8 hours. So a huge story from yesterday will naturally lose the top spot to a smaller story from this morning, without anyone having to manually replace it.
- ↓ **Clickbait and filler get punished.** Anything that looks like a "best deals of the week," "roundup," "newsletter digest," or "how to" guide automatically gets its score cut to a quarter — so those never accidentally become the big headline.

In one sentence: the top story is whichever recent story (posted in the last day and a half) has the highest combination of "sounds important" keywords, source reputation, and freshness — recalculated fresh every time the page loads.

This exact same scoring system is also what decides the order of stories everywhere else on the site — the "Trending Now" row and the order within each topic row are just this same point system, sorted highest to lowest.

5. "What If the Original Story Gets Deleted or Turns Out Fake?"

This is a really good question to be able to answer honestly, so here's exactly what happens:

Honest answer: it does not get removed right away. My platform only checks "is this story still listed on the news site's feed right now?" — it can't tell the difference between "the news site is just having a slow moment and didn't send this story back this time" and "the news site deliberately deleted or retracted this story." So, to be safe, it deliberately keeps showing a story for up to **7 days** after it was first published, even if it stops appearing in the source's feed.

Why build it this way on purpose? Because news websites occasionally go down for a few minutes, or get temporarily blocked, or their feed hiccups. If the platform deleted a story the instant a site failed to list it again, entire sections of the homepage could randomly vanish and reappear every few minutes — which looked broken in earlier versions. So the current design deliberately "assumes innocent" and keeps a story around for a week rather than aggressively deleting.

So what actually happens if a story is later found to be fake?

- If the news site itself publishes a correction or retraction, that correction is just a brand-new article — it would show up on my platform as its own new story, separately.
- The original (now-wrong) story doesn't get automatically pulled down — it simply ages out naturally after 7 days like everything else.
- My platform doesn't do any fact-checking of its own — it trusts whatever the original news site published. If the news site is reliable, the story is reliable; I'm only a window into their reporting, not a verifier of it.

Good talking point: "I don't try to police what's true — I'm a discovery layer on top of trusted publishers. If they're wrong, my job is to always link straight back to them so you're reading their correction too, not a copy frozen on my site."

6. How the Short Summaries Are Written (Right Now, Without AI)

This is the part that's changed. The platform was originally built so an AI (Anthropic's Claude) would read each headline and write a short, neutral 2-3 sentence summary of it. That needs something called an **API key** — basically a paid password that unlocks Anthropic's AI writing service. Right now, that key **is not set up**, so the AI writer is switched off.

What's happening instead — the "extractive" method: when the AI is off, the platform falls back to a simpler, completely free method. Every RSS feed already includes a short preview snippet written by the news site itself (usually the article's opening line or two). My platform takes that snippet, trims it down to a clean, readable length (roughly 50-90 words), and shows that as the "summary." It's not AI-written — it's simply the publisher's own preview text, lightly cleaned up and shortened.

How to explain this to someone in one line

"Right now my summaries are just the news site's own preview text, shortened to a few sentences — not AI-generated. The platform is built to use an AI writer instead, I just haven't switched that on yet."

What changes once the AI key is added later

Nothing about the design changes — the same summary box on the same card just starts being filled in by the AI instead of the shortened preview text. The AI would only ever be given the headline and that same short preview snippet (never the full article), and told to write a neutral 2-3 sentence summary of what happened, who's involved, and why it matters — no opinions, no clickbait phrasing. If the AI ever fails to respond for some reason, the site automatically falls back to the same simple preview-text method described above, so the site never breaks either way.

7. The Whole Thing in 60 Seconds

- 1 Every minute, check 36 established news sites' auto-updating headline lists (RSS feeds) — big tech newsrooms, mainstream papers' tech desks, AI specialists, official company blogs, security researchers, and India-focused outlets.
- 2 Clean up each headline, sort it into a topic, remove duplicates.
- 3 Give every story a score based on how "important-sounding" its words are, how reputable the source is, and how fresh it is.
- 4 The highest-scoring recent story becomes the big top headline; everything else is sorted by score into rows.
- 5 Each story gets a short summary — right now that's just a trimmed version of the news site's own preview text, since the AI writer isn't switched on.
- 6 Old stories quietly fade in ranking and disappear after about a week; nothing is ever manually deleted or fact-checked by the platform itself.
- 7 Every headline links straight back to the real article on the real news site.

Written in plain language from the actual, current behavior of the codebase — not a general description. Reflects the platform exactly as it runs today, including that AI summarization is currently off and the exact 36 sources currently configured.